

Working Title: the data science canon

Databrew

2021-04-13



# Contents

<b>1</b>	<b>Welcome</b>	<b>11</b>
<b>I</b>	<b>Core theory</b>	<b>13</b>
<b>2</b>	<b>Principles of data science</b>	<b>15</b>
2.1	What is data science? . . . . .	15
2.2	What is the data life cycle? . . . . .	15
2.3	What is a pipeline? . . . . .	15
2.4	Data science ‘in the wild’ . . . . .	15
2.5	The reproducibility crisis . . . . .	15
<b>3</b>	<b>Visualizing data</b>	<b>17</b>
3.1	Bad examples . . . . .	17
3.2	Good exaples . . . . .	17
3.3	Edward Tufte . . . . .	17
3.4	Grammar of graphics . . . . .	17
3.5	Design principles . . . . .	17
3.6	Plots & power . . . . .	17
<b>4</b>	<b>Writing about data</b>	<b>19</b>
<b>5</b>	<b>Data ethics</b>	<b>21</b>

<b>II</b>	<b>Getting started</b>	<b>23</b>
<b>6</b>	<b>Setting up RStudio</b>	<b>25</b>
<b>7</b>	<b>Running R code</b>	<b>27</b>
	Learning goals . . . . .	27
	Tutorial video . . . . .	27
	R and RStudio: what's the difference? . . . . .	27
	RStudio's <i>Console</i> . . . . .	28
	Running code in the <i>Console</i> . . . . .	28
	Use R like a calculator . . . . .	30
	7.1 Using operators in R . . . . .	31
	7.2 Use built-in functions within R . . . . .	33
	Review assignment: . . . . .	34
	7.3 Other Resources . . . . .	35
<b>8</b>	<b>Introduction to Rstudio</b>	<b>37</b>
	Learning goals . . . . .	37
	Tutorial video . . . . .	37
	8.1 Tour of RStudio . . . . .	38
	8.2 Other tabs . . . . .	38
	8.3 Scripts . . . . .	38
	8.4 Typical workflows . . . . .	38
	8.5 Other Resources . . . . .	38
<b>9</b>	<b>Variables in R</b>	<b>39</b>
	Learning goals . . . . .	39
	Introducing variables . . . . .	39
	Types of data in R . . . . .	42
	Review assignment . . . . .	43
	Other Resources . . . . .	43

<i>CONTENTS</i>	5
<b>10 Structures for data in R</b>	<b>45</b>
Learning goals . . . . .	45
Introducing data structures . . . . .	45
Vectors . . . . .	45
Review assignment . . . . .	52
Other Resources . . . . .	52
<b>11 Calling functions</b>	<b>53</b>
<b>12 Base plots</b>	<b>55</b>
Learning goals . . . . .	55
Introduction . . . . .	55
Create a basic plot . . . . .	55
Most common types of plots . . . . .	56
Basic plot formatting . . . . .	58
Plotting with data frames . . . . .	66
Next-level plotting . . . . .	67
Review assignment . . . . .	74
Other Resources . . . . .	74
<b>13 Packages</b>	<b>75</b>
<b>14 Basics of ggplot</b>	<b>77</b>
 <b>III Working with data in R</b>	 <b>79</b>
<b>15 Importing data</b>	<b>81</b>
15.1 Working directories . . . . .	81
15.2 Reading in data . . . . .	81
<b>16 Dataframes</b>	<b>83</b>
16.1 Exploration . . . . .	83
16.2 Summarization . . . . .	83

<b>17 Data wrangling</b>	<b>85</b>
17.1 Data transformation . . . . .	85
17.2 The tidyverse and tibbles . . . . .	85
17.3 Transformation with dplyr . . . . .	85
 <b>IV Exploring &amp; analyzing data</b>	 <b>87</b>
<b>18 Exploratory Data Analysis</b>	<b>89</b>
18.1 Exploring distributions . . . . .	89
18.2 Variable types & statistics . . . . .	89
18.3 Descriptive statistics . . . . .	89
 <b>19 Significance statistics</b>	 <b>91</b>
19.1 Thinking about significance . . . . .	91
19.2 Comparison tests . . . . .	91
19.3 Correlation tests . . . . .	91
 <b>20 Displaying data</b>	 <b>93</b>
20.1 Tables . . . . .	93
20.2 Base plots . . . . .	93
20.3 ggplot . . . . .	93
 <b>V Creating your own dataset</b>	 <b>95</b>
<b>21 Managing project files</b>	<b>97</b>
<b>22 Formatting your own data</b>	<b>99</b>
<b>23 Reading Excel files</b>	<b>101</b>
<b>24 Reading GoogleSheets</b>	<b>103</b>
<b>25 Reading online data</b>	<b>105</b>

<i>CONTENTS</i>	7
<b>VI Your R tool bag</b>	<b>107</b>
<b>26 Joining datasets</b>	<b>109</b>
<b>27 for loops</b>	<b>111</b>
Learning goals . . . . .	111
Coming soon . . . . .	111
Tutorial video . . . . .	111
Basics . . . . .	111
Using <code>for</code> loops with data . . . . .	113
Using a <code>for</code> loop with more complex data . . . . .	117
Review assignment . . . . .	122
<b>28 Writing functions</b>	<b>129</b>
Learning goals . . . . .	129
Introduction . . . . .	129
Review assignment . . . . .	129
Other Resources . . . . .	129
<b>29 Working with text</b>	<b>131</b>
<b>30 Working with dates &amp; times</b>	<b>133</b>
<b>31 Working with factors</b>	<b>135</b>
<b>32 Cleaning messy data</b>	<b>137</b>
<b>33 Matrices &amp; lists</b>	<b>139</b>
<b>34 Pipes</b>	<b>141</b>
<b>35 Exporting data &amp; plots</b>	<b>143</b>
<b>VII Interactive dashboards</b>	<b>145</b>
<b>36 Intro to Shiny apps</b>	<b>147</b>

<b>37 Shiny dashboards</b>	<b>149</b>
<b>38 Data entry apps</b>	<b>151</b>
 <b>VIII Databases</b>	 <b>153</b>
<b>39 Introduction</b>	<b>155</b>
39.1 What . . . . .	155
39.2 Why . . . . .	155
39.3 When . . . . .	155
39.4 When not . . . . .	155
 <b>40 Platforms</b>	 <b>157</b>
40.1 PostgreSQL . . . . .	157
40.2 mySQL . . . . .	157
40.3 SQLite . . . . .	157
 <b>41 Alternatives</b>	 <b>159</b>
41.1 NoSQL . . . . .	159
 <b>42 Practices</b>	 <b>161</b>
 <b>IX Documenting your work</b>	 <b>163</b>
 <b>43 R Markdown</b>	 <b>165</b>
 <b>44 Reproducible research</b>	 <b>167</b>
 <b>45 Automated reporting</b>	 <b>169</b>
 <b>46 Formatting standards</b>	 <b>171</b>
46.1 Tables . . . . .	171
46.2 Figures . . . . .	171
46.3 Captions . . . . .	171



<i>CONTENTS</i>	9
<b>X Version control and teamwork</b>	<b>173</b>
47 What is version control?	175
48 What is Git?	177
48.1 Repositories . . . . .	177
48.2 Github . . . . .	177
49 Standard git operations	179
50 A git workflow	181
51 Other git platforms	183
<b>XI Writing about data</b>	<b>185</b>
52 Types of writing	187
52.1 Grant proposals . . . . .	187
52.2 Reports and publications . . . . .	187
52.3 Fundraising . . . . .	187
52.4 Press releases . . . . .	187
53 Elements of style	189
54 Sections of a report	191
54.1 Abstract . . . . .	191
54.2 Introduction . . . . .	191
54.3 Methods . . . . .	191
54.4 Results . . . . .	191
54.5 Discussion . . . . .	191
54.6 Other elements . . . . .	191

<b>XII Creating websites</b>	<b>193</b>
<b>XIII Advanced skills</b>	<b>195</b>
55 Mapping	197
56 Geographic computing & GIS	199
57 Statistical modeling	201
58 Apply family	203
59 Iterative statistics	205
60 Iterative simulations	207
61 Image analysis	209
62 Machine learning	211
63 Template	213
Learning goals . . . . .	213
Tutorial video . . . . .	213
Basics . . . . .	213
Review assignment . . . . .	213
Other Resources . . . . .	214

# Chapter 1

## Welcome

Welcome to *Working Title*, the data science canon by *DataBrew*  
Here is some teacher content.



## Part I

# Core theory



## Chapter 2

# Principles of data science

- 2.1 What is data science?
- 2.2 What is the data life cycle?
- 2.3 What is a pipeline?
- 2.4 Data science ‘in the wild’
- 2.5 The reproducibility crisis





## Chapter 3

# Visualizing data

3.1 Bad examples

3.2 Good examples

3.3 Edward Tufte

3.4 Grammar of graphics

3.5 Design principles

3.6 Plots & power

The politics of graphics



## Chapter 4

# Writing about data



## Chapter 5

# Data ethics



## Part II

# Getting started





## Chapter 6

# Setting up RStudio

**First, download and install R:**

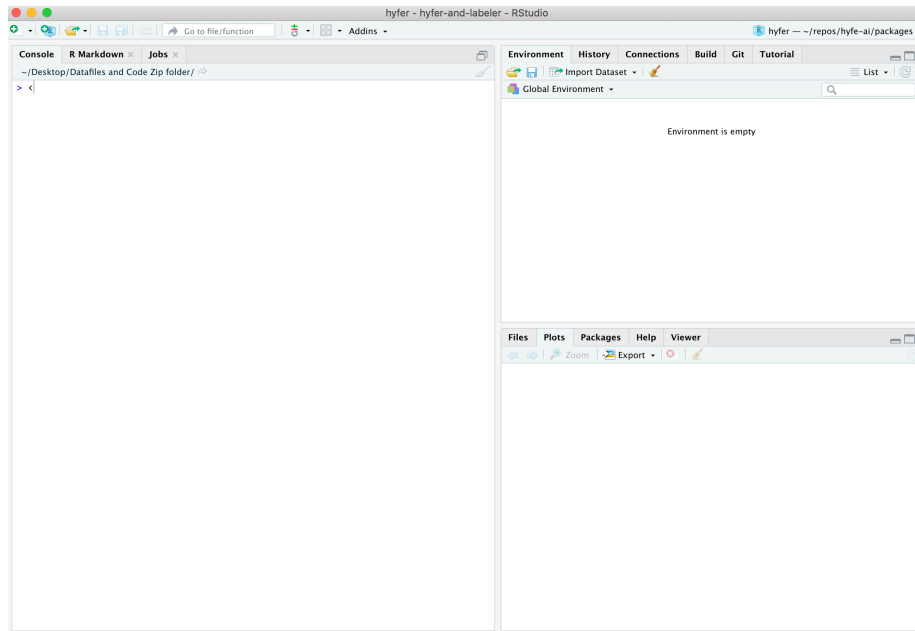
Go to the following website, click the *Download* button, and follow the website's instructions from there. <https://mirrors.nics.utk.edu/cran/>

**Second, download and install RStudio:**

Go to the following website and choose the free Desktop version: <https://rstudio.com/products/rstudio/download/>

**Third, make sure RStudio opens successfully:**

Open the RStudio app. A window should appear that looks like this:



**Fourth, make sure R is running correctly in the background:**

In RStudio, in the pane on the left (the “Console”), type `2+2` and hit Enter. If R is working properly, the number “4” will be printed in the next line down.

**Boom!**

## Chapter 7

# Running R code

### Learning goals

- Understand the difference between R and RStudio.
- Learn how to run code in R
- Learn how to use R as a calculator
- Learn how to use mathematical and logical operators in R

### Tutorial video

#### R and RStudio: what's the difference?

These two things are similar, but it is important to understand how they are different.

In short, R is a open-source (i.e., free) coding language: a powerful programming engine that can be used to do really cool things with data.

**R Studio**, in contrast, is a free *user interface* that helps you interact with **R**. If you think of **R** as an engine, then it helps to think of **RStudio** as the car that contains it. Like a car, **RStudio** makes it easier and more comfortable to use the engine to get where you want to go.

**R Studio** needs **R** in order to function, but **R** can technically be used on its own outside of **RStudio** if you want. However, just as a good car mechanic can get an engine to run without being installed within a car, using **R** on its own takes some expertise. For beginners (and everyone else, really), **R** is just so much more pleasant to use when you are operating it from within **RStudio**.

That is why this book *always* uses **RStudio** when working with **R**.

## RStudio's *Console*

When you open **RStudio**, you see several different panes within the program's window. You will get a tour of **RStudio** in the next module. For now, look at the left half of the screen. You should see a large pane entitled the *Console*.

*NOTE: Insert screenshot here*

**RStudio**'s *Console* is your window into **R**, the engine under the hood. The *Console* is where you type commands for **R** to run, and where **R** prints back the results of what you have told it to do.

## Running code in the *Console*

Type your first command into the *Console*, then press **Enter**:

```
1 + 1
```

```
[1] 2
```

When you press **Enter**, **R** processes the command you fed it, then returns its result (2) just below your command.

Note that spaces don't matter. Both of the following two commands are legible to **R** and return the same thing:

```
4 + 4
```

```
[1] 8
```

```
4+4
```

```
[1] 8
```

However, it is better to make your code as easy to read as possible, which usually means using spaces.

### Exercise 1

Type a command in the *Console* to determine the sum of 596 and 198.

## Re-running code in the *Console*

If you want to re-run the code you just ran, or if you want to recall the code so that you can adjust it slightly, click anywhere in the *Console* then press your keyboard's **Up** arrow.

If you keep pressing your **Up** arrow, R will present you with sequentially older commands.

If you accidentally recalled an old command without meaning to, you can reset the *Console*'s command line by pressing **Escape**.

### Exercise 2

- A. Re-run the sum of 596 and 198 without re-typing it.
- B. Recall the command again, but this time adjust the code to find the sum of 596 and 298.
- C. Practice escaping an accidentally called command: recall your most recent command, then clear the *Console*'s command line.

## Incomplete commands in R

R gets confused when you enter an incomplete command, and will wait for you to write the remainder of your command on the next line in the *Console* before doing anything.

For example, try running this code in your *Console*:

```
45 +
```

You will find that R gives you a little + sign on the line under your command, which means it is waiting for you to complete your command.

If you want to complete your command, add a number (e.g., 3) and hit **Enter**. You should now be given an answer (e.g., 48).

If instead you want R to stop waiting and stop running, hit the **Escape** key.

## Getting errors in R

R only understands your commands if they follow the rules of the R language (often referred to as its *syntax*). If R does not understand your code, it will throw an error and give up on trying to execute that line of code.

For example, try running this code in your *Console*:

```
4 + 6p
```

You probably received a message in red font stating **Error: unexpected symbol in "4 + 6p"**. That is because R did know how to interpret the symbol `p` in this case.

Get used to errors! They happen all the time, even (especially?) to professionals, and it is essential that you get used to reading your own code to find and fix its errors.

### Exercise 3

Type a command in R that throws an error, then recall the command and revise so that R can understand it.

## Use R like a calculator

As you can tell from those commands you just ran, R is, at heart, a fancy calculator.

Some calculations are straightforward, like addition and subtraction:

```
490 + 1000
```

```
[1] 1490
```

```
490 - 1000
```

```
[1] -510
```

Division is pretty straightforward too:

```
24 / 2
```

```
[1] 12
```

For multiplication, use an asterisk (\*):

```
24 * 2
```

```
[1] 48
```

R is usually great about following classic rules for Order of Operations, and you can use parentheses to exert control over that order. For example, these two commands produce different results:

```
2*7 - 2*5 / 2
```

```
[1] 9
```

```
(2*7 - 2*5) / 2
```

```
[1] 2
```

You denote exponents like this:

```
2 ^2
```

```
[1] 4
```

```
2 ^3
```

```
[1] 8
```

```
2 ^4
```

```
[1] 16
```

Finally, note that R is fine with negative numbers:

```
9 + -100
```

```
[1] -91
```

#### Exercise 4

- A. Find the sum of the ages of everyone in your immediate family.
- B. Now recall that command and adjust it to determine the *average* age of the members of your family.

## 7.1 Using operators in R

You can get R to evaluate logical tests using *operators*.

For example, you can ask whether two values are equal to each other.

```
96 == 95
```

```
[1] FALSE
```

```
95 + 2 == 95 + 2
```

```
[1] TRUE
```

R is telling you that the first statement is **FALSE** (96 is not, in fact, equal to 95) and that the second statement is **TRUE** (95 + 2 is, in fact, equal to itself).

Note the use of *double* equal signs here. You must use two of them in order for R to understand that you are asking for this logical test.

You can also ask if two values are *NOT* equal to each other:

```
96 != 95
```

```
[1] TRUE
```

```
95 + 2 != 95 + 2
```

```
[1] FALSE
```

This test is a bit more difficult to understand: In the first statement, R is telling you that it is **TRUE** that 96 is different from 95. In the second statement, R is saying that it is **FALSE** that 95 + 2 is not the same as itself.

Note that R lets you write these tests another, even more confusing way:

```
! 96 == 95
```

```
[1] TRUE
```

```
! 95 + 2 == 95 + 2
```

```
[1] FALSE
```

The first line of code is asking R whether it is not true that 96 and 95 are equal to each other, which is **TRUE**. The second line of code is asking R whether it is not true that 95 + 2 is the same as itself, which is of course **FALSE**.

Other commonly used operators in R include greater than / less than (> and <), and greater/less than or equal to (>= and <=).



```
100 > 100
```

```
[1] FALSE
```

```
100 >= 100
```

```
[1] TRUE
```

### Exercise 5

A. Write and run a line of code that asks whether these two calculations return the same result:

```
2*7 - 2*5 / 2  
(2*7 - 2*5) / 2
```

B. Now write and run a line of code that asks whether the first calculation is larger than the second:

## 7.2 Use built-in functions within R

R has some built-in functions for common calculations, such as finding square roots and logarithms.

```
sqrt(16)
```

```
[1] 4
```

```
log(4)
```

```
[1] 1.386294
```

Note that the function `log()` is the *natural log* function (i.e., the value that  $e$  must be raised to in order to equal 4). To calculate a base-10 logarithm, use `log10()`.

```
log(10)
```

```
[1] 2.302585
```

```
log10(10)
```

```
[1] 1
```

Another handy function is `round()`, for rounding numbers to a specific number of decimal places.

```
100/3
```

```
[1] 33.33333
```

```
round(100/3)
```

```
[1] 33
```

```
round(100/3,digits=1)
```

```
[1] 33.3
```

```
round(100/3,digits=2)
```

```
[1] 33.33
```

```
round(100/3,digits=3)
```

```
[1] 33.333
```

Finally, R also comes with some built-in values, such as  $\pi$ :

```
pi
```

```
[1] 3.141593
```

### Exercise 6

Find the square root of  $\pi$  and round the answer to the 2 decimal places.

## Review assignment:

*NOTE: Under construction!*

## 7.3 Other Resources

Hobbes Primer, Table 1 (Math Operators, pg. 18) and Table 2 (Logical operators, pg. 22)



## Chapter 8


# Introduction to Rstudio

### Learning goals

- Understand the Rstudio working environment and window panes
- Creating and saving scripts
- Customizing your rstudio

### Tutorial video

## 8.1 Tour of RStudio



## 8.2 Other tabs

## 8.3 Scripts

## 8.4 Typical workflows

## 8.5 Other Resources

## Chapter 9

# Variables in R

### Learning goals

- How to define variables and work with them in R
- Learn the various possible classes of data in R

### Introducing variables

So far we have strictly been using R as a calculator, with commands such as:

```
3 + 5
```

```
[1] 8
```

Of course, R can do much, much more than these basic computations. Your first step in uncovering the potential of R is learning how to use **variables**.

In R, a variable is a convenient way of referring to an underlying value. That value can be as simple as a single number (e.g., 6), or as complex as a spreadsheet that is many Gigabytes in size. It may be useful to think of a variable as a cup; just as cups make it easy to hold your coffee and carry it from the kitchen to the couch, variables make it easy to contain and work with data.

### Declaring variables

To assign numbers or other types of data to a variable, you use the `<` and `-` characters to make the arrow symbol `<-`.

```
x <- 3+5
```

As the direction of the `<-` arrow suggests, this command stores the result of `3 + 5` into the variable `x`.

Unlike before, you did not see `8` printed to the *Console*. That is because the result was stored into `x`.

## Calling variables

If you wanted R to tell you what `x` is, just type the variable name into the *Console* and run that command:

```
x
```

```
[1] 8
```

Want to create a variable but also see its value at the same time? Here's a handy trick:

```
x <- 3*12 ; x
```

```
[1] 36
```

The semicolon simulates hitting **Enter**. It says: first run `x <- 3*12`, then run `x`.

You can also update variables.

```
x <- x * 3 ; x
```

```
[1] 108
```

```
x <- x * 3 ; x
```

```
[1] 324
```

You can also add variables together.

```
x <- 8  
y <- 4.5  
x + y
```

```
[1] 12.5
```



## Naming variables

Variables are case-sensitive! If you misspell a variable name, you will confuse R and get an error.

For example, ask R to tell you the value of capital X. The error message will be `Error: object 'X' not found`, which means R looked in its memory for an object (i.e., a variable) named X and could not find one.

You can make variable names as complicated or simple as you want.

```
supercalifragilistic.expialidocious <- 5
supercalifragilistic.expialidocious # still works
```

```
[1] 5
```

Note that periods and underscores can be used in variable names:

```
my.variable <- 5 # periods can be used
my_variable <- 5 # underscores can be used
```

However, hyphens cannot be used since that symbol is used for subtraction.

Also note that variables are case-sensitive! If you name a variable `My_variable`, R will not recognize it if you refer to it as `My_Variable`.

### 9.0.0.1 Naming theory

Naming variables is a bit of an art. The trick is using names that are clear but are not so complicated that typing them is tedious or prone to errors.

Some names need to be avoided, since R uses them for special purposes. For example, `data` should be avoided, as should `mean`, since both are functions built-in to R and R is liable to interpret them as such instead of as a variable containing your data.

Note that R uses a feature called ‘Tab complete’ to help you type variable names. Begin typing a variable name, such as `supercalifragilistic.expialidocious` from the example above, but after the first few letters press the `Tab` key. R will then give you options for auto-completing your word. Press `Tab` again, or `Enter`, to accept the auto-complete. This is a handy way to avoid typos.

### Exercise 1

A. Estimate how many bananas you’ve eaten in your lifetime and store that value in a variable (choose whatever name you wish).

- B. Now estimate how many ice cream sandwiches you've eaten in your lifetime and store that in a different variable.
- C. Now use these variables to calculate your Banana-to-ICS ratio. Store your result in a third variable, then call that variable in the Console to see your ratio.
- D. Who in the class has the highest ratio? Who has the lowest?

## Types of data in R

So far we have been working exclusively with numeric data. But there are many different data types in R. We call these “types” of data **classes**:

- Decimal values like 4.5 are called **numeric** data.
- Natural numbers like 4 are called **integers**. Integers are also numerics.
- Boolean values (TRUE or FALSE) are called **logical** data.
- Text (or string) values are called **character** data.

In order to be combined, data have to be the same class.

R is able to compute the following commands ...

```
x <- 6
y <- 4
x + y
```

```
[1] 10
```

... but not these:

```
x <- 6
y <- "4"
x + y
```

That's because the quotation marks used in naming `y` causes R to interpret `y` as a **character** class.

To see how R is interpreting variables, you can use the `class()` function:

```
x <- 100
class(x)
```

```
[1] "numeric"
```

```
x <- "100"  
class(x)
```

```
[1] "character"
```

```
x <- 100 == 101  
class(x)
```

```
[1] "logical"
```

Another data type to be aware of is **factors**, but we will deal with them later.

### Exercise 3

*NOTE: UNDER CONSTRUCTION!*

### Review assignment

*NOTE: UNDER CONSTRUCTION!*

### Other Resources



## Chapter 10

# Structures for data in R

### Learning goals

- Learn the various structures of data in R
- How to work with vectors in R.

### Introducing data structures

Data belong to different *classes*, as explained in the previous module, and they can be arranged into various **structures**.

So far we have been dealing only with variables that contain a single value, but the real value of R comes from assigning *entire sets* of data to a variable.

### Vectors

The simplest data structure in R is a **vector**. A vector is simply a set of values. A vector can contain only a single value, as we have been working with thus far, or it can contain many millions of values.

### Declaring and using vectors

To build up a vector in R, use the function `c()`, which is short for “concatenate”.

```
x <- c(5,6,7,8)
x
```

```
[1] 5 6 7 8
```

You can use the `c()` function to concatenate two vectors together:

```
x <- c(5,6,7,8)
y <- c(9,10,11,12)
z <- c(x,y)
z
```

```
[1] 5 6 7 8 9 10 11 12
```

You can also use `c()` to add values to a vector:

```
x <- c(5,6,7,8)
x <- c(x,9)
x
```

```
[1] 5 6 7 8 9
```

When two vectors are of the same length, you can do arithmetic with them:

```
x <- c(5,6,7,8)
y <- c(9,10,11,12)
x + y
```

```
[1] 14 16 18 20
```

```
x - y
```

```
[1] -4 -4 -4 -4
```

```
x * y
```

```
[1] 45 60 77 96
```

```
x / y
```

```
[1] 0.5555556 0.6000000 0.6363636 0.6666667
```

You can also put vectors through logical tests:

```
x <- 1:5
4 == x
```

```
[1] FALSE FALSE FALSE TRUE FALSE
```

This command is asking R to tell you whether each element in `x` is equal to 4.

You can create vectors of any data class:

```
x <- c("Ben", "Joe", "Eric")
x
```

```
[1] "Ben" "Joe" "Eric"
```

```
y <- c(TRUE, TRUE, FALSE)
y
```

```
[1] TRUE TRUE FALSE
```

Note that all values within a vector *must* be of the same class (i.e., data type). You can't combine numerics and characters into the same vector. If you did, R would try to convert the numbers to characters. For example:

```
x <- 4
y <- "6"
z <- c(x, y)
z
```

```
[1] "4" "6"
```

## Useful functions for handling vectors

`length()` tells you the number of elements in a vector:

```
x <- c(5,6)
y <- c(9,10,11,12)

length(x)
```

```
[1] 2
```

```
length(y)
```

```
[1] 4
```

The **colon symbol** `:` creates a vector with every integer occurring between a min and max:

```
x <- 1:10
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

**seq()** allows you to build a vector using evenly spaced *sequence* of values between a min and max:

```
seq(0,100,length=11)
```

```
[1] 0 10 20 30 40 50 60 70 80 90 100
```

In this command, you are telling R to give you a sequence of values from 0 to 100, and you want the length of that vector to be 11. R then figures out the spacing required between each value in order to make that happen.

Alternatively, you can prescribe the interval between values instead of the length:

```
seq(0,100,by=7)
```

```
[1] 0 7 14 21 28 35 42 49 56 63 70 77 84 91 98
```

**rep()** allows you to repeat a single value a specified number of times:



```
rep("Hey!",times=5)
```

```
[1] "Hey!" "Hey!" "Hey!" "Hey!" "Hey!"
```

**head()** and **tail()** can be used to retrieve the first 6 or last 6 elements in a vector, respectively.

```
x <- 1:1000
head(x)
```

```
[1] 1 2 3 4 5 6
```

```
tail(x)
```

```
[1] 995 996 997 998 999 1000
```

You can also adjust how many elements to return:

```
head(x,2)
```

```
[1] 1 2
```

```
tail(x,10)
```

```
[1] 991 992 993 994 995 996 997 998 999 1000
```

**sort()** allows you to order a vector from its smallest value to its largest:

```
x <- c(4,8,1,6,9,2,7,5,3)
sort(x)
```

```
[1] 1 2 3 4 5 6 7 8 9
```

**rev()** lets you reverse the order of elements within a vector:

```
x <- c(4,8,1,6,9,2,7,5,3)
rev(x)
```

```
[1] 3 5 7 2 9 6 1 8 4
```

```
rev(sort(x))
```

```
[1] 9 8 7 6 5 4 3 2 1
```

`which()` allows you to ask, “For which elements of a vector is the following statement true?”

```
x <- 1:10  
which(x==4)
```

```
[1] 4
```

If no values within the vector meet the condition, a vector of length zero will be returned:

```
x <- 1:10  
which(x == 11)
```

```
integer(0)
```

`%in%` is a handy operator that allows you to ask whether a value occurs *within* a vector:

```
x <- 1:10  
4 %in% x
```

```
[1] TRUE
```

```
11 %in% x
```

```
[1] FALSE
```

## Exercise 2

*NOTE: UNDER CONSTRUCTION!*

## Subsetting vectors

Since you will eventually be working with vectors that contain thousands of data points, it will be useful to have some tools for *subsetting* them – that is, looking at only a few select elements at a time.

You can subset a vector using square brackets `[ ]`.

```
x <- 50:100
x[10]
```

```
[1] 59
```

This command is asking R to return the 10th element in the vector `x`.

```
x[10:20]
```

```
[1] 59 60 61 62 63 64 65 66 67 68 69
```

This command is asking R to return elements 10:20 in the vector `x`.

### Exercise 3

A. Figure out how to replicate the `head()` function using your new vector subsetting skills.

B. Now replicate the `tail()` function, using those same skills as well as the `length()` function you just learned.

## Dataframes & other data structures

A **vector** is the most basic data structure in R, and the other structures are built out of vectors.

As a data scientist, the most common data structure you will be working with is a **dataframe**, which is essentially a spreadsheet: a dataset with rows and columns, in which each column represents is a vector of the same class of data.

We will explore dataframes in detail later, but here is a sneak peak at what they look like:

```
df <- data.frame(x=300:310,
                 y=600:610)
df
```

```
   x    y
1 300 600
2 301 601
3 302 602
4 303 603
5 304 604
```

```
6 305 605
7 306 606
8 307 607
9 308 608
10 309 609
11 310 610
```

In this command, we used the `data.frame()` function to combine two vectors into a dataframe with two columns named `x` and `y`. R then saved this result in a new variable named `df`. When we call `df`, R shows us the dataframe.

The great thing about dataframes is that they allow you to relate different data types to each other.

```
df <- data.frame(name=c("Ben","Joe","Eric"),
                  height.inches=c(75,73,80))
df
```

```
  name height.inches
1  Ben             75
2  Joe             73
3  Eric            80
```

This dataframe has one column of class `character` and another of class `numeric`.

The two other most common data structures are `matrices` and `lists`, but we will wait on learning about those. For now, focus on becoming comfortable using vectors and dataframes.

### Exercise 3

*NOTE: UNDER CONSTRUCTION!*

## Review assignment

*NOTE: UNDER CONSTRUCTION!*

## Other Resources

## Chapter 11

# Calling functions



## Chapter 12

# Base plots

### Learning goals

- Make basic plots in R
- Basic adjustments to plot formatting

### Introduction

To learn how to plot, let's first create a dataset to work with:

```
country <- c("USA", "Tanzania", "Japan", "Ctr. Africa Rep.", "China", "Norway", "India")
lifespan <- c(79, 65, 84, 53, 77, 82, 69)
gdp <- c(55335, 2875, 38674, 623, 13102, 84500, 6807)
```

These data come from this publicly available database that compares health and economic indices across countries in 2011.

The `lifespan` column presents the average life expectancy for each country.


The `gdp` column presents the average GDP per capita within that country, which is a common index for the income and wealth of average citizens.

Let's see if there is a relationship between life expectancy and income.

### Create a basic plot

The simplest way to make a basic plot in R is to use its built-in `plot()` function:

```
plot(lifespan ~ gdp)
```



This syntax is saying this: plot column `lifespan` as a function of `gdp`. The symbol `~` denotes “*as a function of*”. This frames `lifespan` as a dependent variable (y axis) that is affected by the independent variable (x axis), which in this case is `gdp`.

Note that `R` uses the variable names you provided as the x- and y-axes. You can adjust these labels however you wish (see formatting section below).

You can also produce this exact same plot using the following syntax:

```
plot(y=lifespan, x=gdp)
```

Choose whichever one is most intuitive to you.


## Most common types of plots

The plot above is a **scatter plot**, and is one of the most common types of plots in data science.

You can turn this into a **line plot** by adding a parameter to the `plot()` function:

```
plot(lifespan ~ gdp, type="l")
```






*What a mess!* Rather than connecting these values in the order you might expect, R connects them in the order that they are listed in their source vectors. This is why line plots tend to be more useful in scenarios such as time series, which are inherently ordered.

Another common plot is the **bar plot**, which uses a different R function:

```
barplot(height=lifespan, names.arg=country)
```



In this command, the parameter `height` determines the height of the bars, and `names.arg` provides the labels to place beneath each bar.

Note that one of the bar labels did not plot, because R did not have enough space to print the labels in a pretty way. Rather than create overlapping text, R is programmed to leave some labels out. You will be able to adjust text size on your own (see next section).

There are many more plot types out there, but let's stop here now.

## Exercise 1

Produce a bar plot that shows the GDP for each country.

## Basic plot formatting

You can adjust the default formatting of plots by adding other inputs to your `plot()` command. To understand all the parameters you can adjust, bring up the help page for this function:


```
?plot
```

If multiple help page options are returned, select the *Generic X-Y Plotting* page from the `base` package. This is the plot function that comes built-in to R.

Here we demonstrate just a few of the most common formatting adjustments you are likely to use:

**Set plot range** using `xlim` (for the x axis) and `ylim` (for the y axis):

```
plot(lifespan ~ gdp, xlim=c(0,60000), ylim=c(40,100))
```



In this command, you are defining axis limits using a 2-element vector (i.e., `c(min,max)`).


Note that it can be easier to read your code if you put each input on a new line, like this:

```
plot(lifespan ~ gdp,  
     xlim=c(0,60000),  
     ylim=c(40,100))
```

Make sure each input line within the function ends with a comma, otherwise you R will get confused and throw an error.


**Set dot type** using the input `pch`:

```
plot(lifespan ~ gdp,pch=16)
```



Set **dot color** using the input `col` (the default is `col="black"`)


```
plot(lifespan ~ gdp,col="red")
```



Here is a great resource for color names in R.


Set **dot size** using the input `cex` (the default is `cex=1`):

```
plot(lifespan ~ gdp, cex=3)
```




Set axis labels using the inputs `xlab` and `ylab`:

```
plot(lifespan ~ gdp, xlab="Gross Domestic Product (GDP)", ylab="Life Expectancy")
```




Set axis number size using the input `cex.axis` (the default is `cex.axis=1`):

```
plot(lifespan ~ gdp, cex.axis=.5)
```




Set axis label size using the input `cex.label` (the default is `cex.lab=1`):

```
plot(lifespan ~ gdp, cex.lab=.5)
```



Set plot margins using the function `par(mar=c())` before you call `plot()`:


```
par(mar=c(5,5,0.5,0.5))
plot(lifespan ~ gdp)
```



In this command, the four numbers in the vector used to define `mar` correspond to the margin for the bottom, left, top, and right sides of the plot, respectively.

**Create a multi-pane plot** using the function `par(mfrow=c())` before you call `plot()`:

```
par(mfrow=c(1,2))
plot(lifespan ~ gdp,col="red",pch=16)
plot(lifespan ~ gdp,col="blue",pch=16)
```



In this command, the two numbers in the vector used to define `mfrow` correspond to the number of rows and columns, respectively, on the entire plot. In this case, you have 1 row of plots with two columns.


Note that you will need to reset the number of panes when you are done with your multi-pane plot!

```
par(mfrow=c(1,1))
```

**Plot dots and lines at once using the input type:**

```
par(mfrow=c(1,2))
plot(lifespan ~ gdp, type="b")
plot(lifespan ~ gdp, type="o")
```






```
par(mfrow=c(1,1))
```

Note the two slightly different formats here.

Use a **logarithmic scale** for one or of your axes using the input `log`

```
plot(lifespan ~ gdp, log="x")
```



## Exercise 2

Produce a *beautifully* formatted plot that incorporates **all** of these customization inputs explained above into a multi-paned plot.

## Plotting with data frames

So far in this tutorial we have been using vectors to produce plots. This is nice for learning, but does not represent the real world very well. You will almost always be producing plots using dataframes.


Let's turn these vectors into a dataframe:

```
df <- data.frame(country,lifespan,gdp)
df
```

	country	lifespan	gdp
1	USA	79	55335
2	Tanzania	65	2875
3	Japan	84	38674
4	Ctr. Africa Rep.	53	623
5	China	77	13102
6	Norway	82	84500
7	India	69	6807

To plot data within a dataframe, your `plot()` syntax changes slightly:

```
plot(lifespan ~ gdp, data=df)
```



This syntax is saying this: using the dataframe named `df` as a source, plot column `lifespan` as a function of column `gdp`. The symbol `~` denotes “*as a function of*”. This frames `lifespan` as a dependent variable (y axis) that is affected by the independent variable (x axis), which in this case is `gdp`.

Another way to write this command is as follows:

```
plot(df$lifespan ~ df$gdp)
```

In this command, the `$` symbol is saying, “give me the column in `df` named `lifespan`”. It is a handy way of referring to a column within a dataframe by name. You will learn more about working with dataframes in an upcoming module.

## Exercise 2

- Use the `df` dataframe to produce a bar plot that shows life expectancy for each country.
- Use the `df` dataframe to produce a jumbled line plot of life expectancy as a function of GDP. Reference the `plot()` documentation to figure out how to change the thickness of the line.

## Next-level plotting

The possibilities for data visualization in R are pretty much limitless, and over time you will become fluent in making gorgeous plots. Here are a few common


tools that can take your plots to the next level.

## Adding lines

In some cases it is useful to add reference lines to your plot. For example, what if we wanted to be able to quickly see which countries had life expectancies below 75 years?

You can add a line at `lifespan = 75` using the function `abline()`.


```
plot(df$lifespan ~ df$gdp)
abline(h=75)
```



In this command, the `h` input means “place a horizontal line at this y value.”.

Similarly, you can use `v` to specify vertical lines at certain x values.

```
plot(df$lifespan ~ df$gdp)
abline(v=c(25000,35000),lty=3)
```



Note here that another input, `lty`, was used to change the type of line printed. (Refer to `?abline()` for more details).


#### Exercise 4

Produce a plot of life expectancy as a function of GDP per capita. Then add a line to your plot that indicates which countries have per-capita GDPs that fall below (or above) the average per-capita GDP for the whole dataset. Make your line dashed and color it red.

#### Adding text

Use the `text()` function to add labels to your plot:

```
plot(df$lifespan ~ df$gdp)
text(x=30000,y=75,labels="My text went here")
```



### Exercise 5


Produce a plot of life expectancy as a function of GDP per capita, then label each point by country. Make the labels small and place them *to the right* of their associated dot (Hint: use `?text` for help).

### Highlighting certain data points

It can be helpful to highlight a certain data point (or group of data points) using a different dot size, format, or color.

**To highlight a single data point**, here is one approach you can take: first, plot all points, *then* re-plot the point of interest using the `points()` function:

```
plot(df$lifespan ~ df$gdp)
points(x=df$gdp[5], y=df$lifespan[5], col="red", pch=16, cex=1.5)
```



In this example, we re-plotted the data for the fifth row in the dataframe (in this case, China).

**To highlight a group of data points**, try this approach:

- First, create a vector that will contain the color for each data point.
- Second, determine the color for each data point using a logical test.
- Third, use your vector of colors within your `plot()` command.

For example, let's highlight all countries whose life expectancy is greater than 75.


```
# First
cols <- rep("grey",times=length(lifespan)) # create a vector of colors the length of vector `lifespan`
cols
```

```
[1] "grey" "grey" "grey" "grey" "grey" "grey" "grey"
```

```
# Second
change_these <- which(lifespan > 75)
change_these # these are the elements that we want to highlight
```

```
[1] 1 3 5 6
```

```
cols[change_these] <- "red" # change the color for these elements to a highlight color
# Third
plot(lifespan ~ gdp, pch=16, col=cols, log="x")
```



### Exercise 6

Produce a plot of life expectancy as a function of GDP per capita, in which all countries with GDPs below \$10,000 have larger dots of a different color.

### Building a plot from the ground up

In many applications it can be helpful to have complete control over the way your plot is built. To do so, you can build your plot from the very bottom up in multiple steps.

The steps for building up your own plot are as follows:

1. **Stage a blank canvas:** A plot begins with a blank canvas that covers a certain range of values for x and y. To stage a blank canvas, add this parameters to your `plot()` function: `type="n", axes=FALSE, ann=FALSE, xlim=c(__, __), ylim=c(__, __)`. These commands tell R to plot a blank space, not to print axes, not to print annotations like x- or y-axis labels, and to limit your canvas to a certain coordinate range. Be sure to add numbers to the `xlim()` and `ylim()` commands.



2. **Add your axes**, if you want, using the function `axis()`. The command `axis(1)` prints the x-axis, and `axis(2)` prints the y-axis. This function allows you to define where tick marks occur and other details (see `?axis`).
3. **Add axis titles** using the function `title()`.
4. **Add reference lines**, if you want, using `abline()`. Do this before adding data, since it is usually nice for data points to be superimposed *on top of* your reference lines.
5. **Add your data** using either `points()` or `lines()`.
6. **Add text labels**, if you want, using `text()`.

Here is an example of this process:

```
# 1. Stage a blank canvas
par(mar=c(4.5,4.5,1,1))
plot(1,type="n",axes=FALSE,ann=FALSE,xlim=c(0,100000),ylim=c(40,100))


# 2. Add axes
axis(1,at=c(0,20000,40000,60000,80000,100000),labels=c("$0", "$20", "$40", "$60", "$80", "$100"))
axis(2,at=seq(40,100,by=10),las=2)

# 3. Add axis titles
title(xlab="Gross Domestic Product (GDP, in thousands) per Capita ",cex.lab=.9)
title(ylab="Average Life Expectancy",cex.lab=.9)

# 4. Add reference lines
abline(h=70,v=50000,lty=3,col="grey")

# 5. Add data
points(x=gdp,y=lifespan,pch=16,col="firebrick")

# 6. Add text
text(x=gdp[6],y=lifespan[6],labels="Norway",pos=3)
```



## Review assignment

*NOTE: Under construction!*

## Other Resources

## Chapter 13

# Packages



## Chapter 14

# Basics of ggplot



## Part III

# Working with data in R





## Chapter 15

# Importing data

### 15.1 Working directories

### 15.2 Reading in data



## Chapter 16

# Dataframes

### 16.1 Exploration

### 16.2 Summarization



## Chapter 17

# Data wrangling

### 17.1 Data transformation

#### 17.1.1 Filtering

#### 17.1.2 Grouping

#### 17.1.3 Joining

### 17.2 The tidyverse and tibbles

### 17.3 Transformation with dplyr

#### 17.3.1 Filtering

#### 17.3.2 Grouping

#### 17.3.3 Mutating



## Part IV

# Exploring & analyzing data





## Chapter 18

# Exploratory Data Analysis

18.1 Exploring distributions

18.2 Variable types & statistics

18.3 Descriptive statistics



## Chapter 19

# Significance statistics

19.1 Thinking about significance

19.2 Comparison tests

19.3 Correlation tests



## Chapter 20

# Displaying data

### 20.1 Tables

### 20.2 Base plots

Advanced techniques

### 20.3 ggplot

Advanced techniques



## Part V

# Creating your own dataset





## Chapter 21

# Managing project files



## Chapter 22

# Formatting your own data



## Chapter 23

# Reading Excel files



## Chapter 24

# Reading GoogleSheets





## Chapter 25

# Reading online data



## Part VI

# Your R tool bag



## Chapter 26

# Joining datasets



# Chapter 27

## for loops

### Learning goals

- What `for` loops are, and how to use them yourself
- How to use `for` loops for multi-pane plotting
- How to use `for` loops to achieve complex plots
- How to use `for` loops to summarize data efficiently

### Coming soon

- Instructor notes and answer keys (hidden from students)

### Tutorial video

*(coming soon!)*

### Basics

A `for` loop is a super powerful coding tool. In a `for` loop, `R` loops through a chunk of code for a set number of repetitions.

A super basic example:

```
x <- 1:5
for(i in x){
  print(i)
}
```

```
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
```

Here's an example of a pretty useless for loop:

```
for(i in 1:5){
  print("I'm just repeating myself.")
}
```

```
[1] "I'm just repeating myself."
[1] "I'm just repeating myself."
[1] "I'm just repeating myself."
[1] "I'm just repeating myself."
[1] "I'm just repeating myself."
```

**This code is saying:**

- For each iteration of this loop, step to the next value in `x` (first example) or `1:5` (second example).
- Store that value in an object `i`,
- and run the code inside the curly brackets. - Repeat until the end of `x`.

**Look at the basic structure:**

- In the `for( )` parenthetical, you tell R what values to step through (`x`), and how to refer to the value in each iteration (`i`).
- Within the curly brackets, you place the chunk of code you want to repeat.

Another basic example, demonstrating that you can update a variable repeatedly in a loop.

```
x <- 2
for(i in 1:5){
  x <- x*x
  print(x)
}
```

```
[1] 4
```



```
[1] 16
[1] 256
[1] 65536
[1] 4294967296
```

Another silly example:

```
professors <- c("Keri","Deb","Ken")
for(x in professors){
  print(paste0(x," is pretty cool!"))
}
```

```
[1] "Keri is pretty cool!"
[1] "Deb is pretty cool!"
[1] "Ken is pretty cool!"
```

## Exercise 1

Use this space to practice the basics of `for` loop formatting.

First, create a vector of names (add at least 3)

```
# Add your names to this vector
famous.names <- c("Lady Gaga","David Haskell","Tom Cruise")
```

Using the examples above as a guide, create a `for` loop that prints the same silly statement about each of these names.

```
# Do your coding here
for(i in famous.names){
  print(paste0(i," has cooties!"))
}
```

```
[1] "Lady Gaga has cooties!"
[1] "David Haskell has cooties!"
[1] "Tom Cruise has cooties!"
```

## Using for loops with data

These silly examples above do a poor job of demonstrating how powerful a `for` loop can be.

## Multi-panel plots

For example, a `for` loop can be a very efficient way of making multi-panel plots.

Let's use a `for` loop to get a quick overview of the variables included in the `airquality` dataset built into R.

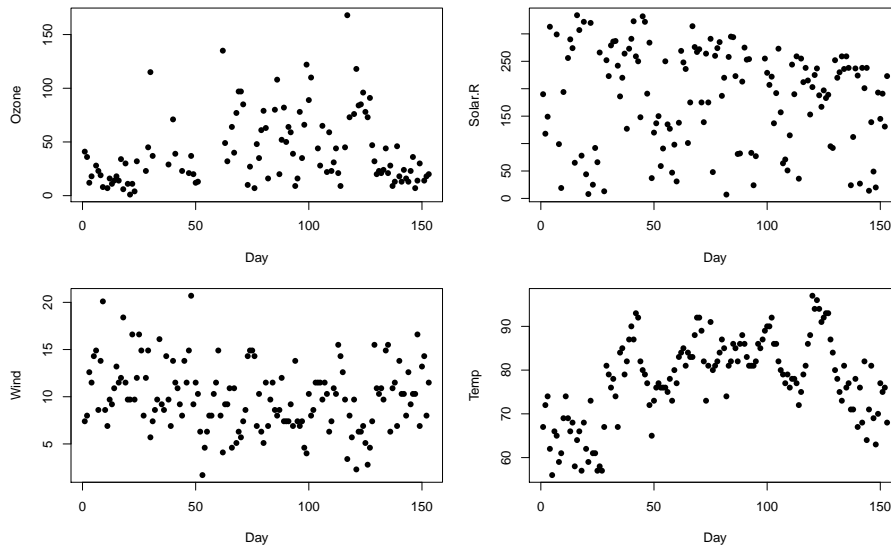
```
data(airquality)
head(airquality)
```

	Ozone	Solar.R	Wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8.0	72	5	2
3	12	149	12.6	74	5	3
4	18	313	11.5	62	5	4
5	NA	NA	14.3	56	5	5
6	28	NA	14.9	66	5	6

Looks like the first four columns would be interesting to plot.

```
par(mfrow=c(2,2)) # Setup a multi-panel plot # format = c(number of rows, number of co
par(mar=c(4.5,4.5,1,1)) # Set plot margins

for(i in 1:4){
  y <- airquality[,i]
  var.name <- names(airquality)[i]
  plot(y,xlab="Day",ylab=var.name,pch=16)
}
```



```
par(mfrow=c(1,1)) # restore the default single-panel plot
```

## Tricky plot solutions

for loops are also useful for plotting data in tricky ways. Let's use a different built-in dataset, that shows the performance of various car make/models.

```
data(mtcars)
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Let's say we want to see how gas mileage is affected by the number of cylinders a car has. It would be nice to create a plot that shows the raw data as well as the mean mileage for each cylinder number.

```
# Let's see how many different cylinder types there are in the data
ucyl <- unique(mtcars$cyl) ; ucyl
```

```
[1] 6 4 8
```


```
# Let's make an empty plot
plot(1,type="n", # tell R not to draw anything
     xlim=c(2,10),ylim=c(0,50),
     xlab="Number of cylinders",
     ylab="Gas mileage (mpg)")

# Write your for loop here to add the actual data
i=ucyl[1] # It's always good to use a known value of i as you build up your for loop
for(i in ucyl){

  # Subset the dataframe according to number of cylinders
  cari <- mtcars[mtcars$cyl==i,]

  # Plot the raw data
  points(x=cari$cyl,y=cari$mpg,col="grey")

  # Superimpose the mean on top
  points(x=i,y=mean(cari$mpg),col="black",pch="-",cex=5,)
}
```



## Exercise 2


Now try to do something similar on your own with the `airquality` dataset. Use `for` loops to create a plot with Month on the x axis and Temperature on

the y axis. On this plot, depict all the temperatures recorded in each month in the color grey, then superimpose the mean temperature for each month.

We will provide the empty plot, you provide the `for` loop:

```
plot(1,type="n",
     xlim=c(3,10),ylim=c(40,100),
     xlab="Month",
     ylab="Yemperature")

# Write your for loop here to add the actual data
for(i in airquality$Month){
  airi <- airquality[airquality$Month==i,]
  points(x=airi$Month,y=airi$Temp,pch=1,col="grey")
  points(x=i,y=mean(airi$Temp),pch="-",cex=5,col="black")
}
```



## Using a `for` loop with more complex data

Here's another good example of the power of a good `for` loop.

First, read in some cool data.

```
kc <- read.csv("../data/keeling-curve.csv") ; head(kc)
```


```
year month day_of_month day_of_year year_dec frac_of_year    CO2
```

1	1974	5	26	145.4890	1974.399	0.3986	332.95
2	1974	6	2	152.4970	1974.418	0.4178	332.35
3	1974	6	9	159.5050	1974.437	0.4370	332.20
4	1974	6	16	166.5130	1974.456	0.4562	332.37
5	1974	6	23	173.4845	1974.475	0.4753	331.73
6	1974	6	30	180.4925	1974.495	0.4945	331.68

This is the famous Keeling Curve dataset: long-term monitoring of atmospheric CO<sub>2</sub> measured at a volcanic observatory in Hawaii.


Try plotting the Keeling Curve:

```
plot(kc$C02 ~ kc$year_dec,type="l",xlab="Year",ylab="Atmospheric C02")
```



There are some erroneous data points! We clearly can't have negative CO<sub>2</sub> values. Let's remove those and try again:

```
kc <- kc[kc$C02 > 0,]
plot(kc$C02 ~ kc$year_dec,type="l",xlab="Year",ylab="Atmospheric C02")
```



**What's the deal with those squiggles?** Let's investigate!

Let's look at the data a different way: *by focusing in on a single year.*

```
# Stage an empty plot for what you are trying to represent
plot(1, # plot a single point
     type="n",
     xlim=c(0,365),xlab="Day of year",
     ylim=c(-5,5),ylab="CO2 anomaly")
abline(h=0,col="grey") # add nifty horizontal line

# Reduce the dataset to a single year (any year)
kcy <- kc[kc$year=="1990",] ; head(kcy)
```

	year	month	day_of_month	day_of_year	year_dec	frac_of_year	CO2
816	1990	1	7	6.4970	1990.018	0.0178	353.58
817	1990	1	14	13.5050	1990.037	0.0370	353.99
818	1990	1	21	20.5130	1990.056	0.0562	353.92
819	1990	1	28	27.4845	1990.075	0.0753	354.39
820	1990	2	4	34.4925	1990.094	0.0945	355.04
821	1990	2	11	41.5005	1990.114	0.1137	355.09

```
# Let's convert each CO2 reading to an 'anomaly' compared to the year's average.
CO2.mean <- mean(kcy$CO2,na.rm=TRUE) ; CO2.mean # Take note of how useful that 'na.rm=TRUE' inpu
```


```
[1] 354.4538
```

```
y <- kcy$CO2 - CO2.mean ; y # Translate each data point to an anomaly
```

```
[1] -0.87384615 -0.46384615 -0.53384615 -0.06384615  0.58615385  0.63615385
[7]  0.96615385  0.72615385  1.13615385  1.33615385  1.08615385  1.67615385
[13]  1.81615385  1.71615385  1.77615385  2.41615385  2.50615385  3.24615385
[19]  2.79615385  2.87615385  2.92615385  2.52615385  1.79615385  1.72615385
[25]  1.33615385  1.76615385  0.53615385 -0.16384615 -0.08384615 -0.46384615
[31] -1.28384615 -0.99384615 -1.37384615 -2.65384615 -3.29384615 -3.59384615
[37] -2.70384615 -2.99384615 -3.05384615 -2.91384615 -2.88384615 -2.72384615
[43] -2.05384615 -1.74384615 -1.30384615 -1.00384615 -0.76384615 -0.55384615
[49]  0.01615385 -0.11384615  0.37615385  0.34615385          NA
```

```
# Add points to your plot
```

```
points(y~kcy$day_of_year,pch=16,col=adjustcolor("darkblue",alpha.f=.3))
```



But this only shows one year of data! How can we include the seasonal squiggle from other years?

Let's use a `for` loop!

OK – let's redo that graph and add a `for` loop into the mix:

```
# First, stage your empty plot:
plot(1,type="n",
     xlim=c(0,365),xlab="Day of year",
     ylim=c(-5,5),ylab="CO2 anomaly")
```




```
abline(h=0,col="grey")
```

```
# Now we will loop through each year of data. First, get a vector of the years included in the d  
years <- unique(kc$year) ; years
```

```
[1] "1974" "1975" "1976" "1977" "1978" "1979" "1980" "1981" "1982" "1983"  
[11] "1984" "1985" "1986" "1987" "1988" "1989" "1990" "1991" "1992" "1993"  
[21] "1994" "1995" "1996" "1997" "1998" "1999" "2000" "2001" "2002" "2003"  
[31] "2004" "2005" "2006" "2007" "2008" "2009" "2010" "2011" "2012" "2013"  
[41] "2014" "2015" "2016" "2017" "2018" "2019" "2020" "2021" NA
```

```
# Now build your for loop.  
# Notice that the contents of the `for loop` are exactly the same  
# as the single plot above -- with one exception.  
# Notice the use of the symbol i  
  
for(i in years){  
  
  # Reduce the dataset to a single year  
  kcy <- kc[kc$year==i,] ; head(kcy)  
  
  # Let's convert each CO2 reading to an 'anomaly' compared to the year's average.  
  CO2.mean <- mean(kcy$CO2,na.rm=TRUE) ; CO2.mean # Get average CO2 for year  
  
  y <- kcy$CO2 - CO2.mean ; y # Translate each data point to an anomaly  
  
  # Add points to your plot  
  points(y~kcy$day_of_year,pch=16,col=adjustcolor("darkblue",alpha.f=.3))  
}
```



Beautiful! So how do you interpret this graph? Why does the squiggle happen every year?

## Review assignment

First, read in and format some other cool data. The code for doing so is provided for you here:

```
df <- read.csv("../data/renewable-energy.csv")
```

This dataset, freely available from World Bank, shows the renewable electricity output for various countries, presented as a percentage of the nation's total electricity output. They provide this data as a time series.

### 27.0.1 Summarize columns with a for loop

**Task 1:** Use a `for` loop to find the change in renewable energy output for each nation in the dataset between 1990 and 2015. Print the difference for each nation in the console.

```
# Write your code here
names(df)
```

```
[1] "year"          "World"         "Australia"     "Canada"
```

```
[5] "China"          "Denmark"        "India"          "Japan"
[9] "New_Zealand"    "Sweden"         "Switzerland"    "United_Kingdom"
[13] "United_States"
```

```
i=2
for(i in 2:ncol(df)){
  dfi <- df[,i] ; dfi
  diffi <- dfi[length(dfi)] - dfi[1] ; diffi
  print(paste0(names(df)[i], " : ",round(diffi,"% change."))
}
```

```
[1] "World : 3% change."
[1] "Australia : 4% change."
[1] "Canada : 1% change."
[1] "China : 4% change."
[1] "Denmark : 62% change."
[1] "India : -9% change."
[1] "Japan : 5% change."
[1] "New_Zealand : 0% change."
[1] "Sweden : 12% change."
[1] "Switzerland : 7% change."
[1] "United_Kingdom : 23% change."
[1] "United_States : 2% change."
```

**Task 2:** Re-do this loop, but instead of printing the differences to the console, save them in a vector.

```
# Write your code here
diffs <- c()
i=2
for(i in 2:ncol(df)){
  dfi <- df[,i] ; dfi
  diffi <- dfi[length(dfi)] - dfi[1] ; diffi
  print(paste0(names(df)[i], " : ",round(diffi,"% change."))
  diffs <- c(diffs,diffi)
}
```

```
[1] "World : 3% change."
[1] "Australia : 4% change."
[1] "Canada : 1% change."
[1] "China : 4% change."
[1] "Denmark : 62% change."
[1] "India : -9% change."
[1] "Japan : 5% change."
```

```
[1] "New_Zealand : 0% change."
[1] "Sweden : 12% change."
[1] "Switzerland : 7% change."
[1] "United_Kingdom : 23% change."
[1] "United_States : 2% change."
```

```
diffs
```

```
[1] 3.49241703 3.98181045 0.63273122 3.51887728 62.33064943 -9.14624362
[7] 4.73004321 0.07524008 12.26263811 7.21543884 23.01128298 1.69994636
```

## Multi-pane plots with for loops

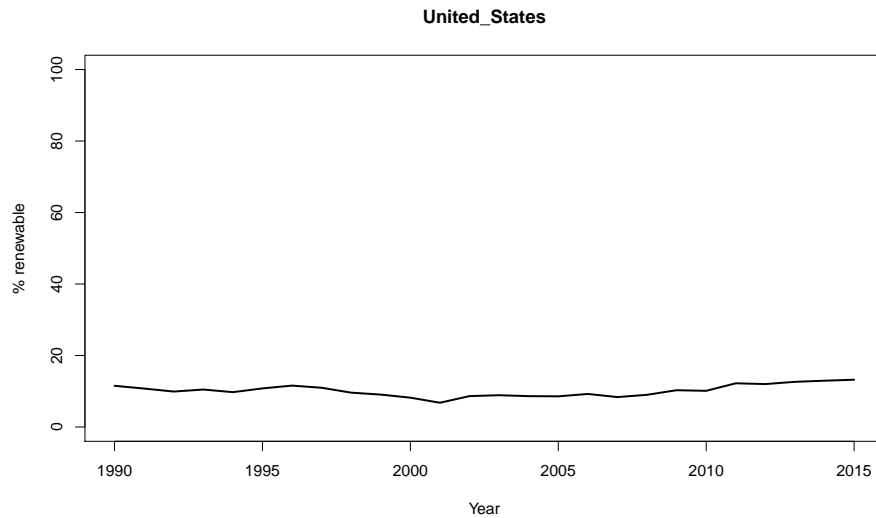
### Practice with a single plot

**Task 3:** First, get your bearings by figuring out how to use the `df` dataset to plot the time series for the United States, for the years 1990 - 2015. Label the x axis “Year” and the y axis “% Renewable”. Include the full name of the county as the main title for the plot.

```
# Write code here
head(df)
```

	year	World	Australia	Canada	China	Denmark	India	Japan
1	1990	19.36204	9.656031	62.37872	20.40794	3.175275	24.48929	11.254738
2	1991	19.23357	10.598201	61.41041	18.47113	2.892325	22.80740	11.856735
3	1992	19.15840	10.066865	61.67921	17.58468	4.398464	20.75265	10.162888
4	1993	19.78795	10.549144	61.72233	18.12526	4.730088	19.55881	11.454528
5	1994	19.53812	10.194474	60.40045	18.08844	4.295431	21.21910	7.993026
6	1995	19.83536	9.624143	61.00410	19.21414	5.035639	17.26054	9.416323
		New_Zealand	Sweden	Switzerland	United_Kingdom	United_States		
1		80.00620	51.00011	54.98254	1.828767	11.528647		
2		77.18945	44.30088	57.16370	1.656439	10.757414		
3		72.58771	52.33321	56.90938	2.005662	9.916110		
4		77.02407	52.92433	59.57279	1.777626	10.484326		
5		82.05216	43.02873	60.57322	2.139842	9.747236		
6		83.85281	47.57878	57.42996	2.066535	10.801085		

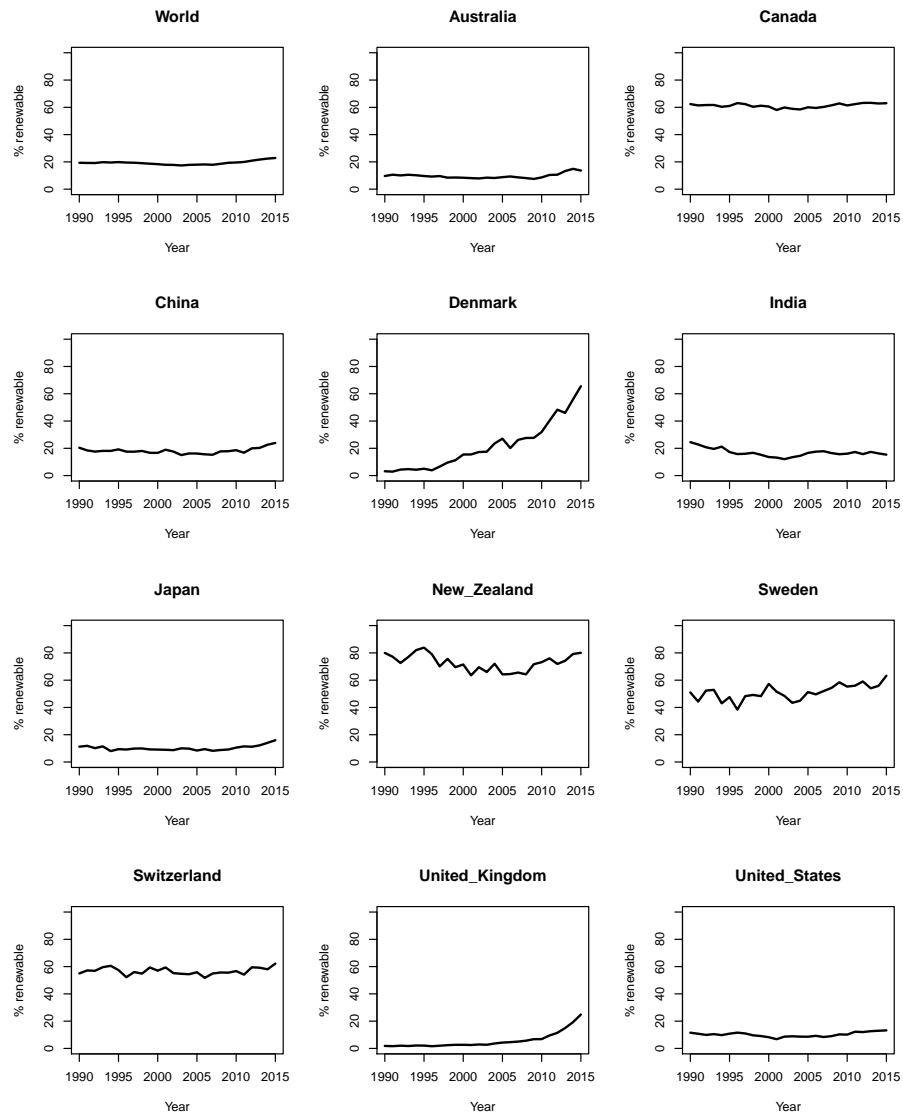
```
dfi <- df[,c(1,13)]
plot(x=dfi[,1],
     y=dfi[,2],
     type="l",lwd=2,
     xlim=c(1990,2015),ylim=c(0,100),
     xlab="Year",ylab="% renewable",
     main=names(dfi)[2])
```



**Now loop it!**

**Task 4:** Use that code as the foundation for building up a `for` loop that displays the same time series for every country in the dataset on a multi-pane graph that with 4 rows and 3 columns.

```
par(mfrow=c(4,3))
i=3
for(i in 2:ncol(df)){
  dfi <- df[,c(1,i)] ; dfi
  plot(x=dfi[,1],
       y=dfi[,2],
       type="l",lwd=2,
       xlim=c(1990,2015),ylim=c(0,100),
       xlab="Year",ylab="% renewable",
       main=names(dfi)[2])
}
```



Now loop it differently!

**Task 5:** Now try a different presentation. Instead of producing 12 different plots, superimpose the time series for each country on the *same single plot*.

To add some flare, highlight the USA curve by coloring it red and making it thicker.

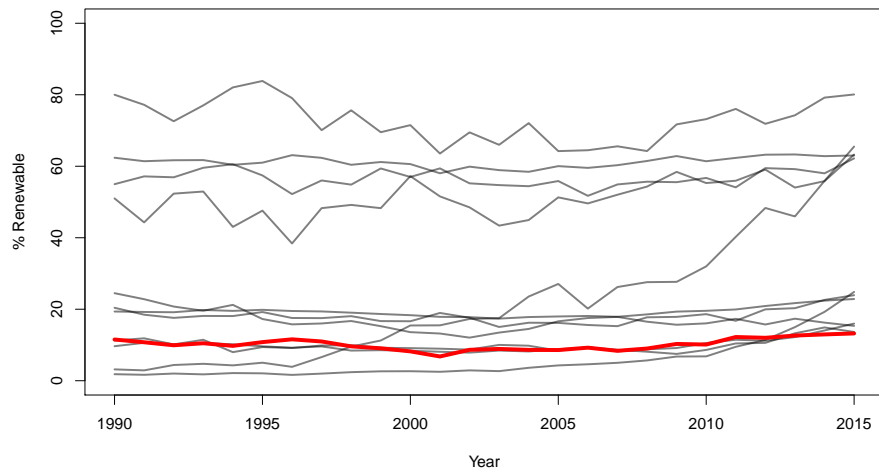
```

par(mfrow=c(1,1))
plot(1,type="n",lwd=2,
     xlim=c(1990,2015),ylim=c(0,100),
     xlab="Year",ylab="% Renewable")

for(i in 2:ncol(df)){
  dfi <- df[,c(1,i)] ; dfi
  lines(dfi[,2]~dfi[,1],lwd=2,col=adjustcolor("black",alpha.f=.5))
}

lines(df$United_States~df$year,lwd=4,col="red")

```







## Chapter 28

# Writing functions

### Learning goals

- Item 1
- Item 2
- Item 3

### Introduction

### Exercise 1

### Review assignment

### Other Resources



## Chapter 29

# Working with text



## Chapter 30

# Working with dates & times



## Chapter 31

# Working with factors





## Chapter 32

# Cleaning messy data



## Chapter 33

# Matrices & lists



## Chapter 34

# Pipes



## Chapter 35

# Exporting data & plots





## Part VII

# Interactive dashboards



## Chapter 36

# Intro to Shiny apps



## Chapter 37

# Shiny dashboards



## Chapter 38

# Data entry apps





**Part VIII**

**Databases**



## Chapter 39

# Introduction

**39.1** What

**39.2** Why

**39.3** When

**39.4** When not



## Chapter 40

# Platforms

40.1 PostgreSQL

40.2 MySQL

40.3 SQLite



## Chapter 41

# Alternatives

### 41.1 NoSQL





## Chapter 42

# Practices

Spinning up a local DB



## Part IX

# Documenting your work



## Chapter 43

# R Markdown



## Chapter 44

# Reproducible research





## Chapter 45

# Automated reporting



## Chapter 46

# Formatting standards

46.1 Tables

46.2 Figures

46.3 Captions



## Part X

# Version control and teamwork



## Chapter 47

# What is version control?





## Chapter 48

# What is Git?

### 48.1 Repositories

### 48.2 Github



## Chapter 49

# Standard git operations



## Chapter 50

# A git workflow



## Chapter 51

### Other git platforms





## Part XI

# Writing about data



## Chapter 52

# Types of writing

52.1 Grant proposals

52.2 Reports and publications

52.3 Fundraising

52.4 Press releases



## Chapter 53

### Elements of style



## Chapter 54

# Sections of a report

54.1 Abstract

54.2 Introduction

54.3 Methods

54.4 Results

54.5 Discussion

54.6 Other elements

54.6.1 Acknowledgments

54.6.2 Literature Cited

54.6.3 Tables

54.6.4 Figures

54.6.5 Supplementary Materials





## Part XII

# Creating websites



## Part XIII

# Advanced skills



## Chapter 55

# Mapping



## Chapter 56

# Geographic computing & GIS





## Chapter 57

# Statistical modeling



## Chapter 58

### Apply family



## Chapter 59

# Iterative statistics



## Chapter 60

# Iterative simulations





## Chapter 61

# Image analysis



## Chapter 62

# Machine learning



# Chapter 63

# Template

## Learning goals

- Item 1
- Item 2
- Item 3

## Tutorial video

Bangarang - Crew Briefing from Luke Padgett on Vimeo.

## Basics

### Exercise 1

## Review assignment

Introduce data

Introduce task(s)

## Other Resources

<https://desiree.rbind.io/post/2020/learnr-iframe/>

<https://rstudio.github.io/learnr/>